

Julia for High Performance Computing

Carsten Bauer

Paderborn Center for Parallel Computing (PC2)

20-23 September 2022

Schedule

	Tuesday	Wednesday	Thursday	Friday
	Foundations	"Serial"	Parallel	Accelerators
9:00- 11:00 (2h)	Introduction Types & Dispatch	Performance Optimisation	Distributed-	Exercises GPU
		Short break		
11:15- 12:45 (1,5h)	Exercises Specialisation &	Exercises Microbench.	Computing Exercises	Programming Q&A
		Lunch break		
14:15- 15:15 (1h)	Generic Prog.	SIMD & LinAlg	Experience Rep.	
		Short break		
15:30 – 16:30 (1h)	Workflow & Pkgs	Profiling	Multithreading	

Course Material

github.com/carstenbauer/JuliaHLRS22

Live Survey

Why Julia for HPC?

The Julia Programming Language

- **Great for domain science**

- “Python-like” syntax
- Dynamic & interactive
 - powerful REPL
 - Jupyter notebooks ([Julia](#) + [Python](#) + [R](#))
- Powerful linear algebra + ecosystem
- Package manager, environments

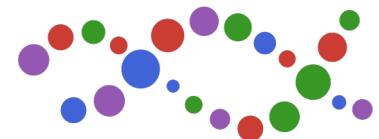
- **Great for HPC**

- Compiled via LLVM (just ahead of time)
- Parallel computing
- Hardware accelerator programming
 - High-level abstractions + low-level kernels
- Interactive HPC

```
julia> function sumup()
           x = 0
           for i in 1:100
               x += i
           end
           return x
       end
sumup (generic function with 1 method)
```

```
julia> @code_llvm debuginfo=:none sumup()
define i64 @julia_sumup_177() #0 {
top:
    ret i64 5050
}
```

Bridging the Gap



Domain Science



julia



HPC

nature

”Julia: come for the syntax, stay for the speed”

[DOI: 10.1038/d41586-019-02310-3](https://doi.org/10.1038/d41586-019-02310-3)

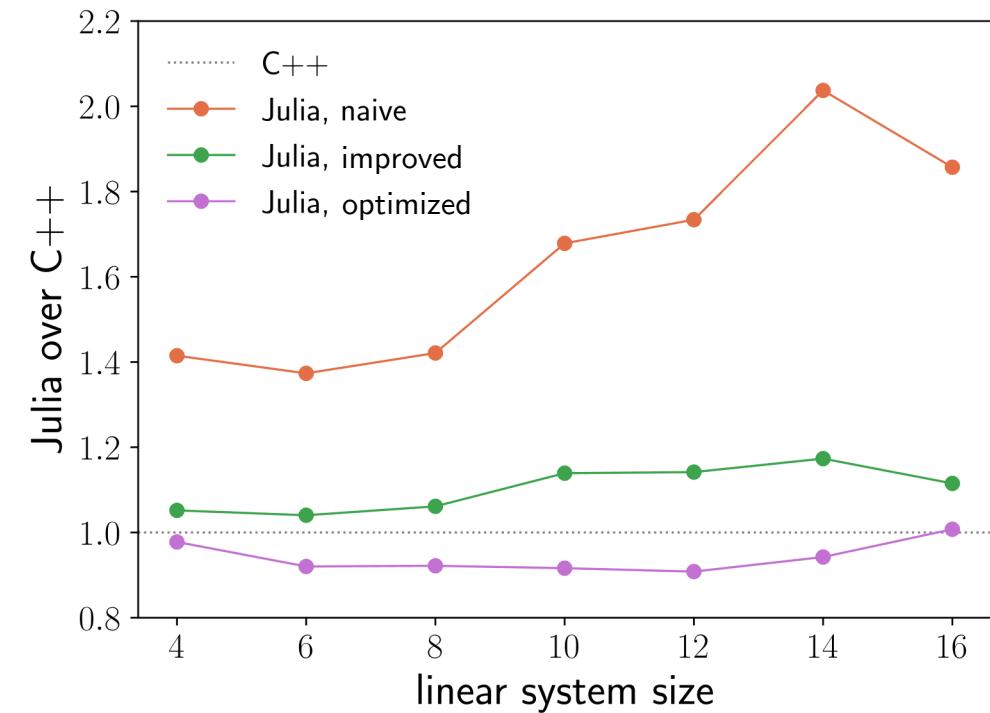
Bridging the Gap



**Julia covers the spectrum between those two “limits”
and can often provide the best of both worlds!**

Personal Experience: Large-Scale Quantum Monte Carlo

- Condensed matter physics application
- Prototypical Julia implementation
 - ~3x faster than Python (not shown)
 - Within a factor of ~2x compared to C++
- Could gradually improve performance!
 - No disruptive language change necessary!
 - A lot of fun!

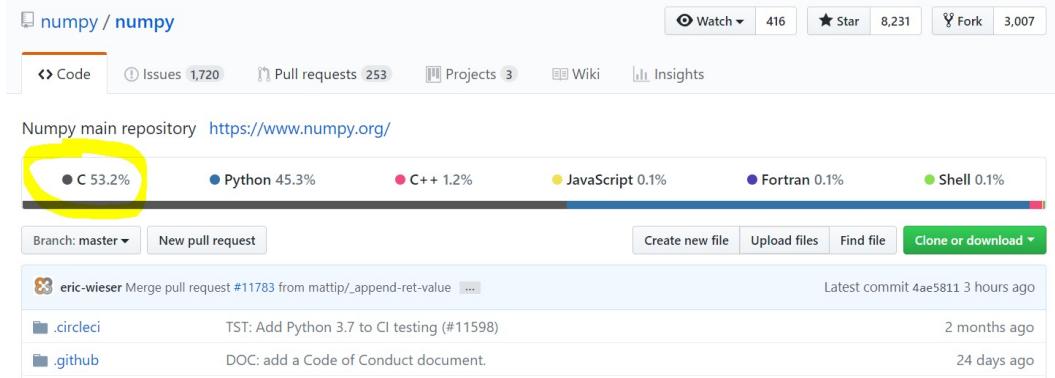


Bridging the Gap



Start with HPC on your laptop and transition to clusters!

Bridging the Gap



Developer



User



Vandermonde matrix

$$V = \begin{bmatrix} 1 & \alpha_1 & \alpha_1^2 & \dots & \alpha_1^{n-1} \\ 1 & \alpha_2 & \alpha_2^2 & \dots & \alpha_2^{n-1} \\ 1 & \alpha_3 & \alpha_3^2 & \dots & \alpha_3^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \alpha_m & \alpha_m^2 & \dots & \alpha_m^{n-1} \end{bmatrix}$$

vander(x)

`numpy.vander(x)`

Python

```
def vander(x, N=None, increasing=False):
    x = asarray(x)
    if x.ndim != 1:
        raise ValueError("x must be a one-dimensional array or sequence.")
    if N is None:
        N = len(x)

    v = empty((len(x), N), dtype=promote_types(x.dtype, int))
    tmp = v[:, ::-1] if not increasing else v

    if N > 0:
        tmp[:, 0] = 1
    if N > 1:
        tmp[:, 1:] = x[:, None]
        multiply.accumulate(tmp[:, 1:], out=tmp[:, 1:], axis=1)

    return v

@TYPE@_kind@(char **args, npy_intp *dimensions, npy_intp *steps, void *NPY_UNUSED(func))
{
    if (IS_BINARY_REDUCE) {
        #if @TYPE@
            @type@ * iop1 = (@type@ *)args[0];
            npy_intp n = dimensions[0];

            *iop1 @OP@= pairwise_sum_@TYPE@(args[1], n, steps[1]);
        #else
            BINARY_REDUCE_LOOP(@type@) {
                io1 @OP@= *(@type@ *)ip2;
            }
            *((@type@ *)iop1) = io1;
        #endif
    }
    else if (!run_binary_simd_@kind@_@TYPE@(args, dimensions, steps)) {
        BINARY_LOOP {
            const @type@ in1 = *(@type@ *)ip1;
            const @type@ in2 = *(@type@ *)ip2;
            *((@type@ *)oip1) = in1 @OP@ in2;
        }
    }
}
```

calls

C code

```
PyFunc_Accumulate(PyFuncObject *ufunc, PyArrayObject *arr, PyArrayObject *out,
                   int axis, int dtype)
{
    PyArrayObject *op[2];
    PyArray_Descr *op_dtypes[2] = {NULL, NULL};
    int op_axes[2] = {op_axes_array[0], op_axes_array[1]};
    npy_uint32 op_flags[2];
    int ndim, ndim, dtype_final;
    int needs_api, need_outer_iterator;

    PyObject *iter = NULL, *iter_inner = NULL;

    /* The selected inner loop */
    PyFuncGenericFunction innerloop = NULL;
    void *innerloopdata = NULL;

    const char *ufunc_name = ufunc_get_name_cstr(ufunc);

    /* These parameters come from extobj or from a TLS global */
    int buffersize = 0, errormask = 0;

    NPY_BEGIN_THREADS_DEF;

    NPY_IF_DBG_PRINT(("Evaluating ufunc %s.accumulate(%s", ufunc_name);

    #if 0
    printf("(Doing %s.accumulate on array with dtype : ", ufunc_name);
    PyObject_Print((PyObject *)PyArray_DESCR(arr), stdout, 0);
    printf("%s)", ufunc_name);
    #endif

    if (_get_buffersize_errormask(NULL, "accumulate", &buffersize, &errormask) < 0) {
        return NULL;
    }

    /* Take a reference to out for later returning */
    Py_XINCREF(out);

    dtype_final = dtype;
    if (get_binary_op_function(ufunc, &otype_final,
                              &innerloop, &innerloopdata) < 0) {
        PyErr_Format(PyExc_ValueError,
                    "could not find a matching type for %s.accumulate.",
                    "requested type has type code '%c'.",
                    ufunc_name, dtype ? dtype->type : '-');

        Py_XDECREF(dtype);
        goto fail;
    }

    if (PyArray_ISALIGNED(arr) || (out && PyArray_ISALIGNED(out))) {
        PyArray_EquivTypes(op_dtypes[0], PyArray_DESCR(arr));
        (out ? PyArray_EquivTypes(op_dtypes[0], PyArray_DESCR(out)) : NULL);
        need_outer_iterator = 1;
    }

    /* If input and output overlap in memory, use iterator to figure it out */
    else if (out != NULL && solve_my_share_memory(out, arr,
                                                    NPY_MAX_SHARE_MEMORY) < 0) {
        need_outer_iterator = 1;
    }

    if (need_outer_iterator) {
        int ndim_iter = 0;
        npy_uint32 flags = NPY_ITER_ZERO_SIZE_OK |
                           NPY_ITER_REF_OK |
                           NPY_ITER_COPY_IF_OVERLAP;
        PyArray_Descr *op_dtypes_param = NULL;

        /* The way accumulate is set up, we can't do buffering,
         * so make a copy instead when necessary.
         */
        ndim_iter = ndim;
        flags |= NPY_ITER_MULTI_INDEX;
    }

    /* Add some more flags.
     *
     * The accumulation outer loop is 'elementwise' over the array, so turn
     * count_m1 = PyArray_DIM(arr);
     */

    /* Get the output */
    if (out == NULL) {
        /* Turn the two items into three for the inner loop */
        dataptr_copy[0] = stridel;
        dataptr_copy[1] = stridel;
        dataptr_copy[2] = stridel;
        NPY_IF_DBG_PRINT(("Iterator loop count %d\n",
                          (int)count_m1));
    }
    else {
        PyArray_Descr *dtype;
        op[0] = out = PyArray_GetOperandArray(iter[0]);
        Py_INCRREF(out);
    }
}

/* Get the variables needed for the loop */
iternext = NPyIter_GetterNext(iter, NULL);
if (iternext == NULL) {
    goto fail;
}

/* Copy the first element to start the reduction.
   *
   * Output (dataptr[0]) and input (dataptr[1]) may point
   * to the same memory, e.g. np.add.accumulate(a, out).
   */
if (dtype == NPY_OBJECT) {
    if (count_m1 == PyArray_DIM(op[0], axis)-1) {
        /* Turn the two items into three for the inner loop */
        dataptr_copy[0] = PyArray_BYTES(op[0]);
        dataptr_copy[1] = PyArray_BYTES(op[1]);
        dataptr_copy[2] = PyArray_BYTES(op[2]);
    }
    else if (iter == NULL) {
        char *dataptr_copy[3];
        npy_intp stride_copy[3];
        int itemsize = op_dtypes[0]->elsize;
        /* Execute the loop with no iterators */
        npy_intp count = PyArray_DIM(op[1], axis);
        npy_intp stride0 = 0, stride1 = PyArray_STRIDE(op[1]),
               stride2 = PyArray_STRIDE(op[2]);
        NPY_IF_DBG_PRINT(("Func: Reduce loop with no
iterators\n"));
        if (PyArray_ISDIN(op[0]) != PyArray_ISDIN(op[1])) {
            PyArray_CompareLists(PyArray_DIMS(op[0]),
                                 PyArray_DIMS(op[1]));
        }
        PyErr_SetString(PyExc_ValueError,
                       "provided out is the wrong size "
                       "for the reduction");
        goto fail;
    }
    npy_intp count_m1, stride, stride;
    stride_copy[0] = stride0;
    stride_copy[1] = stride;
    stride_copy[2] = stride;
    /* Turn the two items into three for the inner loop */
    dataptr_copy[0] = PyArray_BYTES(op[0]);
    dataptr_copy[1] = PyArray_BYTES(op[1]);
    dataptr_copy[2] = PyArray_BYTES(op[2]);
    /* Get the variables needed for the loop */
    iternext = NPyIter_GetterNext(iter, NULL);
    if (iternext == NULL) {
        goto fail;
    }
    /* Copy the first element to start the reduction.
       *
       * Output (dataptr[0]) and input (dataptr[1]) may point
       * to the same memory, e.g. np.add.accumulate(a, out).
       */
    if (dtype == NPY_OBJECT) {
        if (count_m1 == PyArray_DIM(op[0], axis)-1) {
            /* Turn the two items into three for the inner loop */
            dataptr_copy[0] = PyArray_BYTES(op[0]);
            dataptr_copy[1] = PyArray_BYTES(op[1]);
            dataptr_copy[2] = PyArray_BYTES(op[2]);
        }
        else if (iter == NULL) {
            char *dataptr_copy[3];
            npy_intp stride_copy[3];
            int itemsize = op_dtypes[0]->elsize;
            /* Execute the loop with just the outer iterator */
            count_m1 = PyArray_DIM(op[1], axis)-1;
        }
    }
}
```

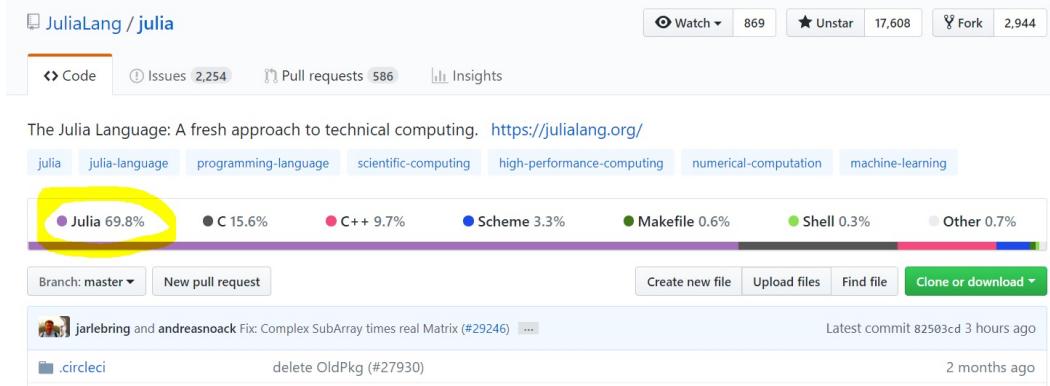
uses

C template

Julia

```
function vander(x::AbstractVector{T}) where T
    m = length(x)
    V = Matrix{T}(undef, m, m)
    for j = 1:m
        V[j,1] = one(x[j])
    end
    for i= 2:m
        for j = 1:m
            V[j,i] = x[j] * V[j,i-1]
        end
    end
    return V
end
```

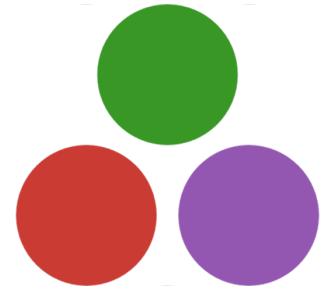
Bridging the Gap



User



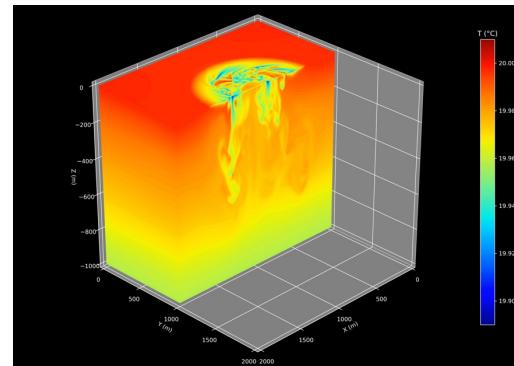
Developer



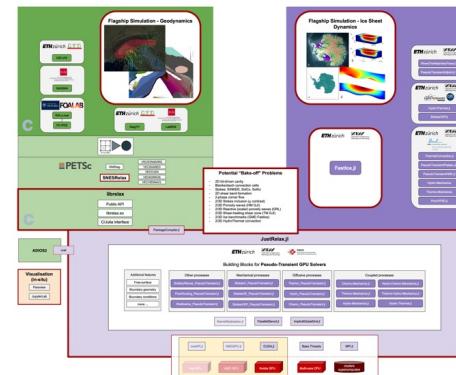
Julia for HPC

International Initiatives

- CliMA @ Caltech, MIT
 - Climate Modeling Alliance
- CESMIX @ MIT
 - Center for the exascale simulation of materials in extreme environment
- ExaSGD @ LLNL, ORNL, ...
 - National power grid optimization
- GPU4GEO @ ETH / CSCS
 - Computational earth science
- ...



Oceananigans.jl



Julia HPC Position Paper

(in preparation)

USA / Canada

- **MIT** Massachusetts Institute of Technology
- **NERSC** National Energy Research Scientific Computing
- **ORNL** Oak Ridge National Laboratory
- **PI** Perimeter Institute for Theoretical Physics

Europe

- **PC2** Paderborn Center for Parallel Computing
- **HLRS** High-Performance Computing Center Stuttgart
- **UHH** University of Hamburg
- **CSCS** Swiss National Supercomputing Centre
- **ETH** Swiss Federal Institute of Technology Zürich
- **ARC@UCL** University College London

Bridging HPC Communities through the Julia Programming Language

Journal Title
XX(X):1-14
OT To Be Assigned (s) 2022
Reprints:
sagepub.co.uk/journalsPermissions.nav
DOI: 10.1177/TobBeAssigned
www.sagepub.com/
SAGE

Valentin Churavy¹, William F Godoy², Carsten Bauer³, Hendrik Ranocha⁴, Michael Schlotke-Lakemper⁵, Ludovic Räss^{6,7}, Johannes Blaschke⁸, Mosè Giordano⁹, Erik Schnetter^{10,11,12}, Samuel Omlin¹³, Jeffrey S. Vetter², Alan Edelman¹

Abstract

The Julia programming language has evolved into a modern alternative to fill existing gaps in the requirements of scientific computing and data science applications. Julia's single-language paradigm, and its proven track record at achieving high-performance without sacrificing user productivity, makes it a viable single-language alternative to the existing composition of high-performance computing (HPC) languages (Fortran, C, C++) and higher-level languages (Python, R, Matlab) suitable for data analysis and simulation alike. Julia's rapid growth in language capabilities, package ecosystem, and community make it a promising new universal language for HPC similar to C++ or Python – an achievable goal if the community is given the necessary resources. This paper presents the views of a multidisciplinary group of researchers in academia, government, and industry advocating for the use of Julia and its ecosystem in HPC centers. We examine the current practice and role of Julia as a common programming model to address major challenges in scientific reproducibility, data-driven artificial intelligence/machine learning (AI/ML), co-design, and in-situ workflows, scalability and performance portability in heterogeneous computing, network, data management, and community education. As a result, we consider necessary the diversification of current investments to fulfill the needs of the upcoming decade as more supercomputing centers prepare for the Exascale era.

Keywords

High Performance Computing, Julia, Programming Language

1 Programming Languages and HPC

The development of programming languages for high performance computing (HPC) has a rich and varied history. Early on, the needs of HPC and main-stream computing were still mostly aligned, leading to the development of Fortran as the first high-level language. Today, we see a bifurcation of programming languages designed for HPC, like Chapel ([Chamberlain et al. 2007](#)), and generalist programming languages, like Python and C++, that are being retooled for HPC use cases. Meanwhile, we see many new languages and programming models being deployed to help users manage the increased heterogeneity of contemporary systems: OpenCL, CUDA, HIP, OpenMP, OpenACC, SYCL, Kokkos, and many others.

Specialist languages lead to a smaller community and make it harder for attracting and developing users who were educated in a more generalist language. We identify these smaller communities as a distinct risk for a language to grow to a sustainable size and to keep attracting users and investments. As an example of the three DARPA HPCS languages Chapel, X10, and Fortress only Chapel is still going.

In generalist languages, the needs of the HPC community are often not addressed or not deemed as important as other concerns. This leads to the development of HPC dialects, based around frameworks.

We will present Julia, a modern generalist programming language, that from the beginning was designed for

¹ Massachusetts Institute of Technology, USA

² Oak Ridge National Laboratory, USA

³ Paderborn Center for Parallel Computing, Paderborn University, Germany

⁴ Department of Mathematics, University of Hamburg, Germany

⁵ High-Performance Computing Center Stuttgart (HLRS), University of Stuttgart, Germany

⁶ Laboratory of Hydraulics, Hydrology and Glaciology (VAW), ETH Zurich, Switzerland

⁷ Swiss Federal Institute for Forest, Snow and Landscape Research (WSL), Birmensdorf, Switzerland

⁸ National Energy Research Scientific Computing Center, Lawrence Berkeley National Laboratory, 1 Cyclotron Road, Berkeley, CA 94720, USA

⁹ Centre for Advanced Research Computing, United Kingdom

¹⁰ Perimeter Institute, 31 Caroline St. N., Waterloo, ON, Canada N2L 2Y5

¹¹ Department of Physics and Astronomy, University of Waterloo, 200 University Avenue West, Waterloo, Ontario, Canada N2L 3G1

¹² Center for Computation & Technology, Louisiana State University, Baton Rouge, LA 70803, USA

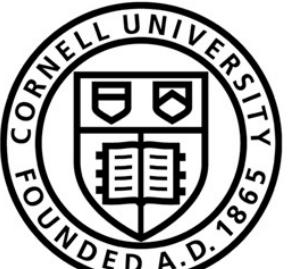
¹³ Swiss National Supercomputing Centre (CSCS), ETH Zurich, Switzerland

Corresponding author:
Valentin Churavy, Computer Science and Artificial Intelligence Laboratory, Massachusetts Institute of Technology, Cambridge, Massachusetts, USA
Email: vchuravy@mit.edu

Julia HPC Community

- Julia Slack
 - Channel: #hpc
 - julialang.org/slack/
- Monthly JuliaHPC Call
 - julialang.org/community/#events
- JuliaCon 2022
 - 27 – 29 July
 - Recording: <https://www.youtube.com/watch?v=fog1x9rs71Q>





EMORY
UNIVERSITY



S
Stanford
University



UNIVERSITY
of
GLASGOW



UNIVERSITE
PAUL
SABATIER



TOULOUSE III



CU THE CITY
UNIVERSITY
OF
NEW YORK

EPFL
ÉCOLE POLYTECHNIQUE
FÉDÉRALE DE LAUSANNE



MŰEGYETEM 1782

TOKYO METROPOLITAN UNIVERSITY
首都大学東京

UCLA

AGH

Let's get started! ☺

github.com/carstenbauer/JuliaHLRS22

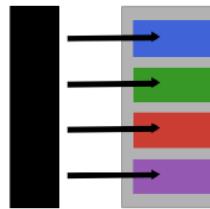
Attribution / Copyright

- Laptop: [Laptop Vectors by Vecteezy](#)
- Cluster: [Web Vectors by Vecteezy](#)

HPC Ecosystem

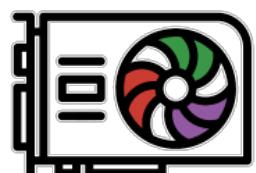
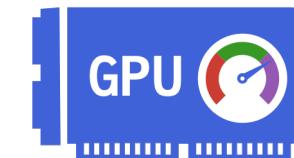
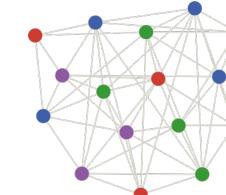
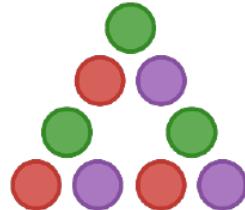
HPC Ecosystem

- JuliaSIMD
 - LoopVectorization.jl, Polyester.jl
- JuliaParallel
 - Dagger.jl, MPI.jl
- JuliaFolds
 - Transducers.jl, FLoops.jl
- JuliaPerf
 - LIKWID.jl, PProf.jl
- JuliaGPU
 - CUDA.jl, AMDGPU.jl



LIKWID.jl

a Julia wrapper for LIKWID

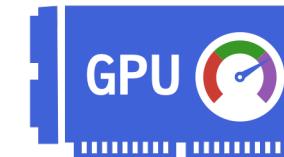
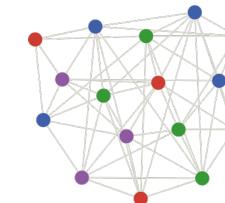


Julia Packages by PC2

- ThreadPinning.jl
 - Dynamically pin Julia threads to CPU threads
- STREAMBenchmark.jl & BandwidthBenchmark.jl
 - Memory bandwidth measurements
- LIKWID.jl
 - Hardware performance monitoring
- GPUInspector.jl
 - Benchmark max. GPU performance
 - GPU node stresstest: monitor temperature & power



LIKWID.jl
a Julia wrapper for LIKWID



pc2.de/go/jlpkg

